

# BEAST Lab Organization

## LRZ

Dr. Josef Weidendorfer, josef.weidendorfer@lrz.de

Amir Raoofy, amir.raoofy@tum.de

## LMU


Sergej Breiter, sergej.breiter@nm.ifi.lmu.de

Minh Thanh Chung, minh.thanh.chung@ifi.lmu.de

Dr. Karl Fűrlinger, karl.fuerlinger@ifi.lmu.de

## TUM

Bengisu Elis, bengisu.elis@tum.de





# Table of Contents

Course Organization

Introduction to BEAST

Assignment 0: Introduction

# Tentative Course Structure

- 14 meetings in total (via Zoom and in class)
  - Last meeting on 9th February 2022
  - Dec 1 - TUM Dies Academicus
- 9 Assignments ( Assignment 0 : optional )
  - 1 week each
- 1 Project
  - 2 weeks
- Student groups of 3 BA students and 2 MA students
- Two groups will present their reports at the meeting right after deadline
  - Presentation notification will be sent 2 days before presentations
  - No slides, go through your report and talk about your findings

# Repository Structure

Gitlab main repository: <https://gitlab.lrz.de/beastlab22ws>

- Each team has a repository, which includes:
  - Lecture slides
  - Assignment material
  - Code template
  - Your code submissions and report (once you place it there)
- Only solutions on the `main` branch will be graded !
  - At the due date, your current master branch state is automatically tagged and archived
  - Make sure your code and report is there by the deadline
- Groups are already set.
- Machine account information will be sent via e-mail - Change passwords please !

# Infrastructure Usage

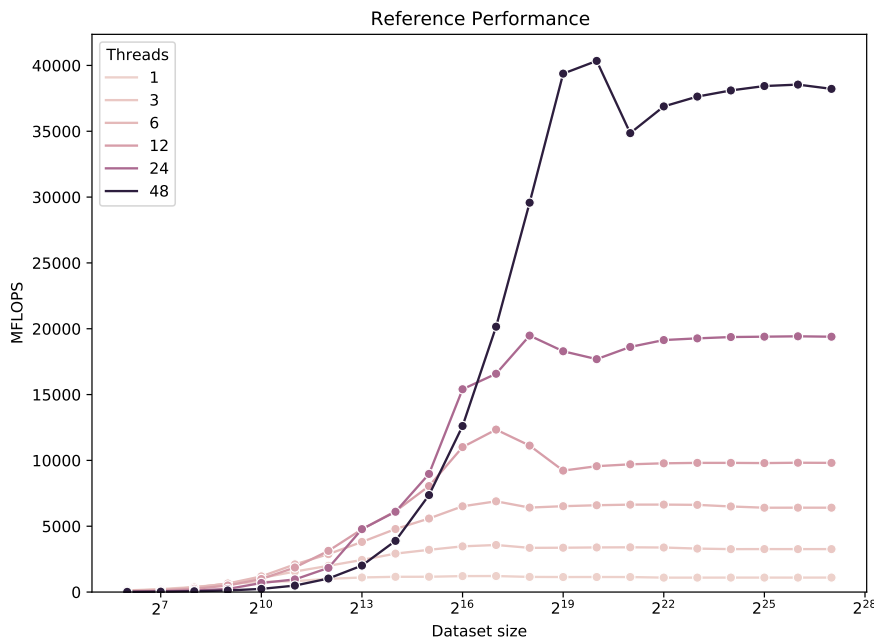
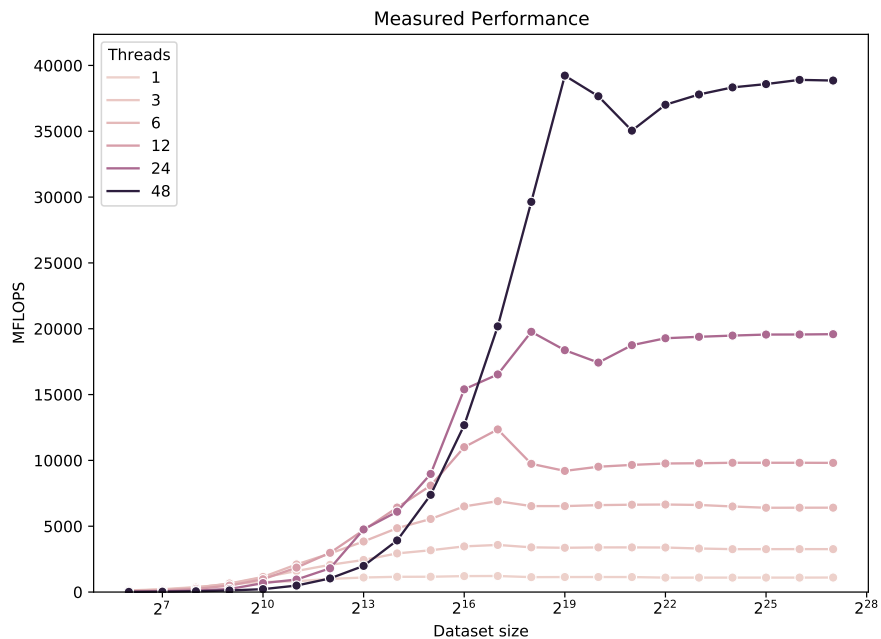
## GitLab

- We need you to sign your commits.
- Tutorial: [https://docs.gitlab.com/ee/user/project/repository/gpg\\_signed\\_commits/](https://docs.gitlab.com/ee/user/project/repository/gpg_signed_commits/)
- If GitLab shows a **Verified** label on your commits you are good to go.
- We will tag your last commit before each deadline.

## Continuous Integration

- Allows for rapid feedback to your code solutions, helps you get set up.
- Code runs on Fujitsu machines for CPU code . Performance data is automatically collected and visualized.

# Sample CI Visualization



Performance data is stored in the build job artifacts (right hand column when viewing build job).

# CI Usage

Benefit: Shows you when your code fails to produce correct results or performance.

	#462227		P master  6e46d7b9 Fix artifact paths for assig...		00:01:07 2 hours ago	
	#462226		P master  ad0dae44 Update artifact paths for s...		00:01:05 2 hours ago	
	#462197		P master  5392ec90 Add lecture slides		00:01:06 2 hours ago	

CI ensures you have the measuring basics correct. You still need to conduct most of the experiments and explain your findings.

## Up Next: Introduction to BEAST



# BEAST Lab

## Introduction to Systems

April 28, 2022 | Josef Weidendorfer

## Collaboration among 3 institutions

LMU  
TUM  
LRZ

LMU – MNM/Prof. Kranzlmüller  
(Karl Förlinger, Minh Chung, Sergej Breiter)

TUM – CAPS/Prof. Schulz  
(Bengisu Elis, Dennis-Florian Herr, Vincent Bode)

LRZ - Future Computing Group  
(Josef Weidendorfer, Amir Raoofy)

We want you to learn about **performance properties of current architectures**

- Be able to understand and explain performance effects seen from measurements
- Get a deeper understanding of current system designs (CPU / GPU)

Part 1: get started with small codes across systems

- We show key hardware design concepts + a parallel programming model (OpenMP)
- We give you typical small HPC code examples
- You run measurements of different scenarios across systems, compare / discuss results
- We all discuss results in weekly meetings, starting with presentations of groups

Structure:

Memory on CPU (Triad / Traversal) → Compute on CPU (MM) → ... on GPU → Tools

We want you to learn about **performance properties of current architectures**

- Be able to understand and explain performance effects seen from measurements
- Get a deeper understanding of current system designs (CPU / GPU)

Part 2: make use of gained knowledge

- We assign randomly one system to each group
- We give you some larger typical HPC code examples
- You tune the code to get best single-node performance (3 week time)
- We discuss intermediate/final experiences/results in weekly meetings

## Target Architectures for the Lab

### CPUs

- Intel Icelake (ISA: x86-64 + AVX512)
- AMD Rome (ISA: x86-64 + AVX2)
- Marvell ThunderX2 (ISA: ARM AArch64 + Neon)
- Fujitsu A64FX (ISA: ARM AArch64 + SVE)

### GPUs

- NVidia V100
- AMD MI-100



SuperMUC-NG  
Top500 (Nov 2018): #8  
**Lenovo Intel (2019)**  
**311,040 cores**  
Intel Xeon Skylake  
**26.9 PetaFlops** Peak  
**19.5 PetaFlops** Linpack  
**719 TeraByte** Main Memory  
**70 PetaByte** Disk

# BEAST – Bavarian Energy Architecture and Software Testbed

## The LRZ Future Computing Testbed



- Help decide about next large system
  - Get experience on benefits of various future architectures for LRZ codes
  - Find best configuration: how much money to spend on compute / memory / network?
  - Enable migration planning: educate own staff / port LRZ tools / prepare courses
  - Support vendor collaboration
- Enable research studies on new technologies
  - Forward looking: LRZ services around future platforms, novel usage models
    - more experimental: FPGAs, AI accelerators, integration of heterogeneity (QC)
  - In partnership with selected researchers from Munich universities

**Lot of work to do! Engage students for student work (BA, MA): This Lab!**



# The Testbed – Available Hardware

2 racks, each with 6 PDUs (for power measurements)

- Max power consumption per rack: 35 kW

Top to bottom (picture from last year)

- 3 switches (Infiniband 200Gb/s HDR), 2x 48port 1Gb/s Ethernet
- Login 1U “testbed.cos.lrz.de”
- **2x AMD Rome** GPU server 2U: “rome1” / “**rome2**”
- Storage 2U with homes
- **2x Marvell ThunderX2** GPU server 2U: “thx1” / “**thx2**”

Not shown:

- HPC CS500 Management 2U + **8 nodes A64FX** “cs1” – “cs8”
- **2x Intel IceLake** GPU server 2U: “ice1” / “**ice2**”



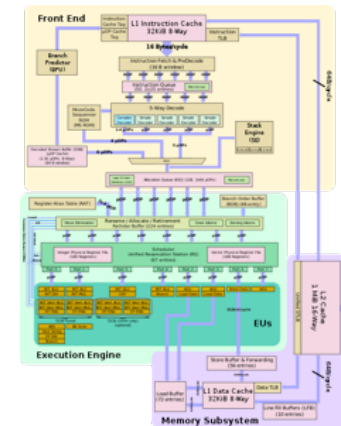
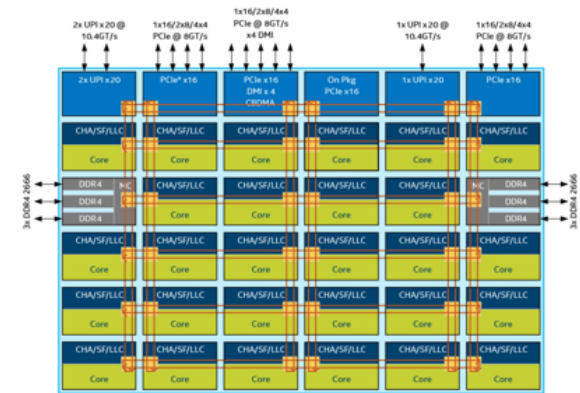
# Intel Skylake (available as fallback)

## SuperMUC-NG

- Here: only Single-Node experiments
- Node
  - 2 sockets with Intel Skylake Xeon Platinum 8174
  - 2x 24 = 48 cores
    - 2x 512bit vector units per core (8 x DP FMA)
    - 2 threads per core ("Hyper-Threading")
    - 2.3 GHz base (currently: 2.5 GHz), 14nm
- 96 GB main memory

## Links

- <https://doku.lrz.de/display/PUBLIC/Hardware+of+SuperMUC-NG>
- [https://en.wikichip.org/wiki/intel/microarchitectures/skylake\\_\(server\)](https://en.wikichip.org/wiki/intel/microarchitectures/skylake_(server))



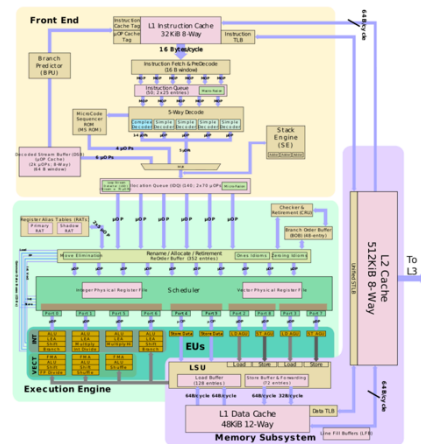
# Intel Icelake

## Two systems in BEAST

- 2 sockets Intel Xeon (Icelake) Platinum 8360Y
  - 2x 36 = 72 cores
    - 2x 512bit vector units per core (8 x DP FMA)
    - 2 threads per core ("Hyper-Threading")
    - 2.4 GHz base, Intel 10nm
- 512 GB main memory, 1.5 TB Optane NVRam

## Links

- [https://en.wikichip.org/wiki/intel/microarchitectures/ice\\_lake\\_\(server\)](https://en.wikichip.org/wiki/intel/microarchitectures/ice_lake_(server))
- [https://en.wikichip.org/wiki/intel/microarchitectures/sunny\\_cove](https://en.wikichip.org/wiki/intel/microarchitectures/sunny_cove)

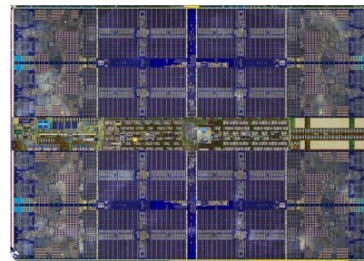


## Two systems in BEAST

- 2 sockets with EPYC 7742
- 2x 64 = 128 cores (“Zen2”)
  - Chiplet design: IO-Die + 8x CCX-Dies (2x 4-core)
  - 2x 256-bit vector units per core (4 x DP FMA)
  - 2 threads per core
  - 2.25 GHz base, TSMC 7nm
- 512 GB main memory
- 2x AMD Radeon MI-100 GPUs
  - 7nm, 32GB HBM, PCIe4

## Link

- [https://en.wikichip.org/wiki/amd/microarchitectures/zen\\_2](https://en.wikichip.org/wiki/amd/microarchitectures/zen_2)



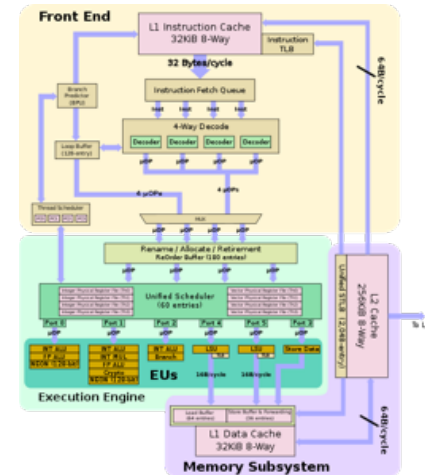
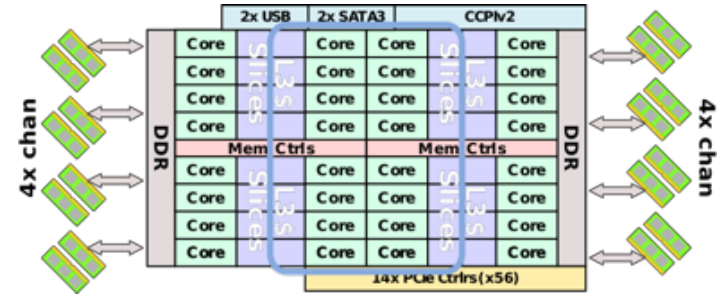
# Marvell ThunderX2

## Two systems in BEAST

- 2 sockets with ThunderX2 CN9980
- 2x 32 = 64 cores (“Vulcan”)
  - 128-bit vector units (2 x DP FMA)
  - 4 threads per core
  - 2.2 GHz base, 16nm
- 512 GB main memory
- 2x Nvidia V-100
  - Volta, 32GB HBM, PCIe3

## Link

- <https://en.wikichip.org/wiki/cavium/microarchitectures/vulcan>



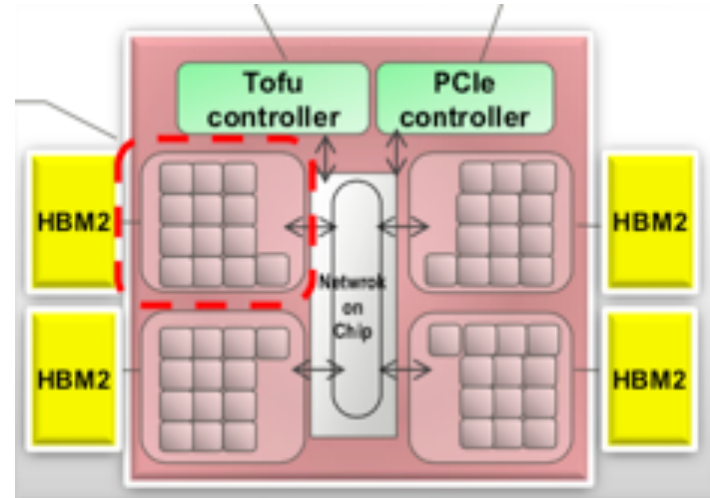
# Fujitsu A64FX

## HPE CS500 in BEAST

- 8 nodes with one A64FX CPU (“NSP1”)
- 48 cores per CPU
  - 2x 512bit vector units per core
  - 1.8 GHz, TSMC 7nm
  - 4 NUMA domains
- 32 GB HBM2

## Link

- [https://en.wikipedia.org/wiki/Fujitsu\\_A64FX](https://en.wikipedia.org/wiki/Fujitsu_A64FX)



[ Fujitsu: The 1<sup>st</sup> SVE Enabled Arm Processor: A64FX and Building up ARM HPC Ecosystem, 2019 ]

## Access

- via “ssh rbgaccount@lxxhalle.informatik.tu-muenchen.de”
- ssh XXX@skx.supermuc.lrz.de

## Compilers

- via “module”, see available packages with “module avail” / load with “module load”
- ICC (“icpc” for C++, “icc” for C, Intel compiler, enable OpenMP: “-openmp”)
- GCC (“g++” C++, “gcc” C, GNU Compiler Collection, enable OpenMP: “-fopenmp”)

## Node allocation (test queue, 30 minutes, e.g. project h039y - see “groups”)

- salloc -Ah039y -ptest -t30

## Access via Linux Cluster login nodes

- `ssh XXX@lxlogin8.lrz.de` (or `lxlogin1@lrz.de`)
- `ssh testbed.cos.lrz.de`
- `ssh <system>`

If `testbed.cos.lrz.de` is not reachable, retry after 1 hour

- probably just a reboot

## Compilers

- system: “gcc”
- via modules: see “module avail”, then “module load <package>”



# Access and Usage: Intel Icelake @ BEAST



## Access

- `ssh XXX@lxlogin8.lrz.de` (or `lxlogin1@lrz.de`)
- `ssh testbed.cos.lrz.de`
- `ssh ice2`

## Compilers

- `gcc`, `icc` (Intel compiler)

# Access and Usage: AMD Rome @ BEAST



## Access

- `ssh XXX@lxlogin8.lrz.de` (or `lxlogin1@lrz.de`)
- `ssh testbed.cos.lrz.de`
- `ssh rome2`

## Compilers

- `gcc`, `clang` (from AMD RocM)

# Access and Usage: ThunderX2 @ BEAST



## Access

- `ssh XXX@lxlogin1.lrz.de`
- `ssh testbed.cos.lrz.de`
- `ssh thx2`

## Compilers

- `gcc`
- (via „module load cuda/11.1.1 llvm“) `clang`

# Access and Usage: AMD A64FX @ BEAST



## Access

- `ssh XXX@lxlogin1.lrz.de`
- `ssh testbed.cos.lrz.de`
- `ssh cs1 / cs2`

## Compilers

- `gcc (8)`
- `gcc 11` (via „`module load gcc/11.0.0`“)
- Cray compiler: “`cc`”, enable OpenMP: “`-h omp`”



Leibniz Supercomputing Centre  
of the Bavarian Academy of Sciences and Humanities



## Up Next: Assignment 0: Introduction